

**Exercise 1:** Here are some data where the observations are divided into an  $x$  group and a  $y$  group.

```
data_mat <- structure(c(4, 5, 9, 10, 6, 4, 13, 4, 16, 14, 9, 2, 7, 6,
9, 16, 7, 7, 7, 9, 1, 15, 11, 20, 8, 5, 19, 16, 17, 18, 18, 17,
20, 15, 29, 16, 14, 17, 17, 25, 17, 1, 12, 33, 26, 16, 3, 25, 7, 2,
30, 16, 24, 6, 3, 23, 24, 24, 15, 2.02, 3.54, 2.27, 4.97, 4.2, 3.15,
4.07, 3.3, 4.2, 4.55, 5.3, 3.86, 4.44, 5.65, 3.23, 6.4, 3.61, 4.63,
4.34, 3.58, 2.13, 5.1, 4.43, 5, 2.46, 2.77, 3.67, 5.67, 5.26, 5.03,
3.25, 5.21, 2.85, 3.87, 4.57, 4.79, 4.17, 4.91, 3.5, 3.42, 3.86, 3.45,
4.73, 2.77, 2.76, 2.67, 4.09, 3.46, 4.03, 3.23, 2.64, 5.92, 4.54, 3.55,
3.43, 4.2, 3.88, 3.59, 3.48, 5.66, 2.48, 6.57, 7.52, 11.47, 6.26, 4.83,
7.6, 9.54, 11.35, 9.96, 10, 7.4, 5.58, 8.91, 8.64, 10.77, 6.04, 9.28, 8.47,
5.72), .Dim = c(20L, 7L), .Dimnames = list(NULL, c("y1", "y2", "y3", "x1", "x2", "x3", "x4")))
```

- Using R, run the Mantel test to see if the distance structure in  $y_1, y_2, y_3$  (count data) is the same as the distance structure for the  $x_1, x_2, x_3, x_4$  data.

**Solution:**

Before we can compute a Mantel test, we must first obtain a distance measure for each group. In the examples that were shown in class we standardized each variable then computed the euclidean distance between each observation (We will also test the count data with the Bray-Curtis measure using the `vegan` package). Proceeding similarly we get a p-value of 0.0002, so at the  $\alpha = .05$  level we would reject the null and conclude that the distance measures for each group are correlated, ie they cluster in a similar way. Plotting the null distribution, and double checking our mantel test with the `mantel()` function from the `vegan` package we get the same conclusion,

**Code:**

```
## Scaling each variable to mean = 0 sd = 1
data_mat_scale <- scale(data_mat)

## Computing the distance matrices for each group.
Y <- dist(data_mat_scale[,1:3], method="euclidean", diag=TRUE, upper=TRUE)
X <- dist(data_mat_scale[,4:7], method="euclidean", diag=TRUE, upper=TRUE)

## Computing the test statistic for the our given data.
T <- cor(cbind(as.vector(Y), as.vector(X)))[1,2]
[1] 0.468862

## Computing the mantel test, generating approx.
## null dist. histogram and computing test statistic.
Nsim <- 5000
null_dist <- rep(NA, Nsim)
Y <- data_mat[,1:3]

for(i in 1:Nsim){
```

```

#Finding a permutation of the 20 observations in dat_mat_scale
ind <- sample(1:20, size=20, replace=FALSE)
#Applying permutation to Y
Y_rand <- Y[ind,]
#Computing permuted distance
Y_rand_dist <- dist(Y_rand, method="", diag=TRUE, upper=TRUE)
#Computing test statistic with permuted Y dist matrix
null_dist[i] = cor(cbind(as.vector(X), as.vector(Y_rand_dist)))[1,2]
}

```

```

hist(null_dist, n=40, xlim=c(-0.4,1.0)) # Plotting null dist.
## Computing P-value
p_value <- mean((null_dist < -abs(T))|(null_dist > abs(T)))
      [1] 2e-04

```

```

## Double checking with mantel test in Vegan library
## Also using the Bray-Curtis distance for the count data
## instead.
load(vegan)
Ymantel <- vegdist(data_mat[,1:3], method = 'bray')
Xmantel <- vegdist(data_mat_scale[,4:7], method = 'euclidean')

mantel(Xmantel, Ymantel, method = 'pearson', permutations = 5000)
Mantel statistic based on Pearson's product-moment correlation


```

```

Call:
mantel(xdis = Xmantel, ydis = Ymantel, method = "pearson", permutations = 5000)
Mantel statistic r: 0.4251
      Significance: 0.0009998
Upper quantiles of permutations (null model):
      90%    95%  97.5%    99%
0.164 0.210 0.254 0.303
Permutation: free
Number of permutations: 5000

```

Figure 1: Distribution of test statistic under the null.



- b. What are your conclusions? what does the Mantel test actually test? What is the null hypothesis?

**Solution:**

Like we previously stated, we conclude that the groups have correlated distance measures we can also say that the groups cluster in a similar way. The null of the Mantel test assumes that the distance measures are uncorrelated, the alternative states that they are correlated. When we do a randomization test, we sample the distribution of the test statistic under the assumption of the null hypothesis since distances should be uncorrelated when the order of observations in one of the groups is randomized. With that sample distribution and our original test statistic we can compute a p-value.

**Exercise 2:** The following is a phylogenetic distance matrix. Use multidimensional scaling (metric, that is, `cmdscale`) to make a map of these species. What does it mean?

HomoSapiens	0.000	0.089	0.104	0.161	0.182	0.232	0.233	0.249	0.256	0.273	0.322	0.308
Pan	0.089	0.000	0.106	0.171	0.189	0.243	0.251	0.268	0.249	0.284	0.321	0.309
Gorilla	0.104	0.106	0.000	0.166	0.189	0.237	0.235	0.262	0.244	0.271	0.314	0.293
Pongo	0.161	0.171	0.166	0.000	0.188	0.244	0.247	0.262	0.241	0.284	0.303	0.293
Hylobates	0.182	0.189	0.189	0.188	0.000	0.247	0.239	0.257	0.242	0.269	0.309	0.296
MacacaFuscata	0.232	0.243	0.237	0.244	0.247	0.000	0.036	0.084	0.124	0.289	0.314	0.282
MacacaMulatta	0.233	0.251	0.235	0.247	0.239	0.036	0.000	0.093	0.120	0.293	0.316	0.289
MacacaFascicular	0.249	0.268	0.262	0.262	0.257	0.084	0.093	0.000	0.123	0.287	0.311	0.298
MacacaSylvanus	0.256	0.249	0.244	0.241	0.242	0.124	0.120	0.123	0.000	0.287	0.319	0.287
SaimiriSciureus	0.273	0.284	0.271	0.284	0.269	0.289	0.293	0.287	0.287	0.000	0.320	0.285
TarsiusSyrichtha	0.322	0.321	0.314	0.303	0.309	0.314	0.316	0.311	0.319	0.320	0.000	0.252
LemurCatta	0.308	0.309	0.293	0.293	0.296	0.282	0.289	0.298	0.287	0.285	0.252	0.000

**Solution:**

Running the multidimensional scaling function `cmd scale` we get the following plot of each of the observations.

**Code:**

```
CMD <- cmdscale(as.dist(D), k = 2)


CMDGOF <- cmdscale(as.dist(D), list=TRUE, k = 2)$GOF
[1] 0.5939294 0.5942756

cmdscale(as.dist(D), list=TRUE, k = 3)$GOF
[1] 0.7328829 0.7333100

plot(MDS, pch=19, xlim = c(-.20, .20), ylim = c(.10, -.25))
text(MDS[,1] -.03, MDS[,2] -.02, rownames(D), cex=.6)

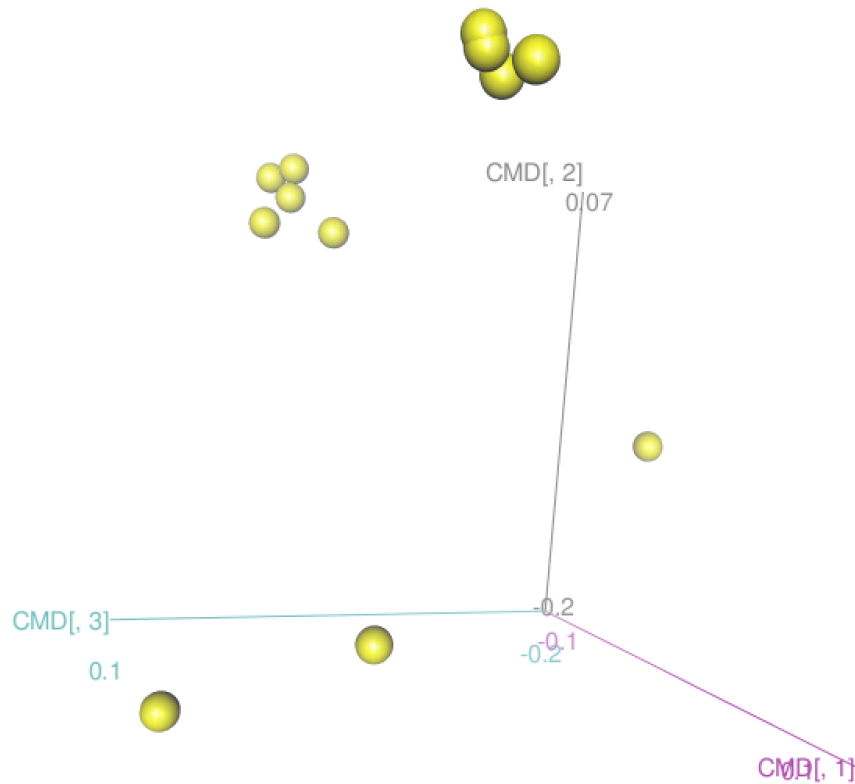
scatter3d(x = CMD[,1], y = CMD[,2], z = CMD[,3], surface = FALSE)
```

Figure 2: MDS plot of phylogenetic distance matrix.



We can see some obvious clustering going on in 2-dimensions. The bottom left corner contains larger primates like Gorilla, Homosapians, Pan and Pongo. the bottom right we can see a clustering of different species in the Macaca genus. The upper cluster seem to be smaller primates like Lemurs and Trasius. When we look at the GOF of our multidimensional scaling we see that we kept around 59% of the variance after compressing the dimension. It might be worthwhile to explore visualizing the data in three dimensions, because when we look at the GOF in three dimensions we get around 73% of the variance. Doing so we get the following plot,

Figure 3: MDS plot of phylogenetic distance matrix with 3 dimensions.



This shows us that the lower group of observations isn't really a cluster. It's possible that in higher dimensions there is a greater distance among the group of large primates, I would assume because of the relation in genus the other cluster stays relatively close together in higher dimensions.

**Exercise 3:** Using the R base function `princomp()` does principal component analysis. Run the following.

```
X <- data_mat[,4:7]
```

```
summary(X)
pairs(X)
#
#
tmp <- princomp(X)
plot(tmp)
tmp
summary(tmp)
#
tmp2 <- princomp(X, cor=TRUE)
plot(tmp2)tmp2summary(tmp2)
```

- a. Does using `cor=TRUE` make a difference? In general, when should you not use `cor=TRUE`? When should you use `cor=TRUE`? Try multiplying column `x1` by 100 and run `princomp()` both ways. Does this support your thoughts about correlation vs covariance matrices in principal components?

### Solution:

Yes, using the correlation matrix for principal component analysis removes the effect of scale for all of the variables, essentially standardizing each variable. If the scale matters in our analysis then we would want to use the covariance matrix. If scale doesn't matter then the correlation matrix is suitable and even better conditioned for the eigenvalue problem. Beyond that if you have a constant variable, the variance will be 0 and the related correlations will be undefined.

Running the following code we see this in action. When we scale the first variable by 100 our covariance matrix PCA attributes all the variance in the data to that first variable. While the correlation matrix PCA is invariant to change in scale.

### Code:

```
-----
#### PCA with X and Covariance Matrix
summary(princomp(X))
Importance of components:
      Comp.1      Comp.2      Comp.3      Comp.4
Standard deviation  2.5105558  1.0157902  0.79133114  0.43315887
Proportion of Variance 0.7734982  0.1266274  0.07684863  0.02302576
Cumulative Proportion 0.7734982  0.9001256  0.97697424  1.00000000
-----
#### PCA with X and Correlation Matrix
summary(princomp(X, cor = TRUE))
Importance of components:
      Comp.1      Comp.2      Comp.3      Comp.4
Standard deviation  1.4408041  1.0741502  0.8147347  0.32633153
Proportion of Variance 0.5189791  0.2884497  0.1659481  0.02662307
Cumulative Proportion 0.5189791  0.8074288  0.9733769  1.00000000
-----
#### Scaling the first variable by 100
```

```
XScaledUP <- cbind(X[,1]*100, X[, (2:4)])

#### PCA with Scaled X and Covariance Matrix
summary(princomp(XScaledUP))
Importance of components:
```

	Comp.1	Comp.2	Comp.3	Comp.4
Standard deviation	103.4506348	1.9300487517	9.157441e-01	4.780835e-01
Proportion of Variance	0.9995524	0.0003479171	7.832267e-05	2.134749e-05
Cumulative Proportion	0.9995524	0.9999003298	9.999787e-01	1.000000e+00

```
-----
#### PCA with Scaled X and Correlation Matrix
summary(princomp(XScaledUP, cor = TRUE))
Importance of components:
```

	Comp.1	Comp.2	Comp.3	Comp.4
Standard deviation	1.4408041	1.0741502	0.8147347	0.32633153
Proportion of Variance	0.5189791	0.2884497	0.1659481	0.02662307
Cumulative Proportion	0.5189791	0.8074288	0.9733769	1.00000000

```
-----
```

- b. Looking at the variances (for instance, when `cor=TRUE`), how many principal components do you think you should use? Why? Does it look like principal components was a good approach to follow in this case, why or why not?

### Solution:

It seems as though there are several rule of thumb test to decide how many components we should retain. I've hear we include them from largest to smallest until we've explained 80% of the variance. I've heard the same rule for 90% and I've also heard that you should drop any components that contribute less than 10% to the variance. I think that looking at our data it seems as though we only need to retain 3 of the 4 components.

I like to think of the components as the magnitude of the semi-axis of the hyperellipse which bounds the data. When we have a small value for a component there exists a hyperellipse in a smaller dimension which is a good approximation for the larger dimension. Generally when we have small components in our data, we are adding extra unneeded complexity(both computationally, and parsimoniously) in the form of another dimension. Whenever we can effectively reduce the dimension of our data with PCA, it is worthwhile.

- c. What does the following set of loadings and scores tell you (in general, what are loadings and scores and what are they used for)?

```
X <- data_mat[,4:7]
hold = princomp(X, cor=TRUE, scores=TRUE)
names(hold)
```



```
hold$scores
hold$loadings
```

**Solution:**

Running the given code we get the following loadings and scores.

**Code:**

```
> hold$scores
      Comp.1      Comp.2      Comp.3      Comp.4
[1,] -3.59837512  0.46592358 -0.3087188 -0.04902597
[2,] -0.10501162 -0.32382941 -0.9637492  0.81770366
[3,] -1.03858755 -1.50611264 -0.7762246 -0.57497050
[4,]  2.13004445  0.68264826 -0.5335500 -0.12447069
[5,] -1.02326106  1.85187203  0.2233541 -0.04203984
[6,] -1.80092901  1.50535704 -0.6359655  0.19873641
[7,] -0.35930163 -0.04003031  0.3851416 -0.17866740
[8,]  0.92084766 -0.72621391 -1.4723743  0.06481649
[9,]  1.59340467 -0.81261185 -0.3868939 -0.50560097
[10,] 1.40325871  0.14930224 -0.4962171  0.12098110
[11,] 1.03431067  1.81051677  0.4264689 -0.39770549
[12,] -0.12086602 -2.56378551  0.7111325  0.21298274
[13,] -1.25216137  0.09466807  1.3447405  0.09458907
[14,]  1.02000282  0.76562254  0.9965436  0.21105188
[15,]  0.08973018 -0.14030215 -1.0270956 -0.22983631
[16,]  2.26714084 -0.12109288  1.4461451  0.23233223
[17,] -0.76243448 -0.21277685 -0.2243883  0.44816422
[18,]  1.11665484 -0.08876730 -0.1424104  0.23257586
[19,]  0.06495144  0.66935439  0.2354269 -0.31574799
[20,] -1.57941842 -1.45974212  1.1986344 -0.21586849
```

```
> hold$loadings
Loadings:
      Comp.1 Comp.2 Comp.3 Comp.4
x1  0.520  0.281  0.706  0.389
x2  0.500 -0.527 -0.447  0.523
x3 -0.168 -0.802  0.544 -0.183
x4  0.672          -0.736
```

```
      Comp.1 Comp.2 Comp.3 Comp.4
SS loadings      1.00  1.00  1.00  1.00
Proportion Var   0.25  0.25  0.25  0.25
Cumulative Var   0.25  0.50  0.75  1.00
```

The loadings tell us how much of each variable should be weighted when computing the component, and the scores tell us the weight of each component in a given observation.

- d. Repeat a (correlation-based) principal component analysis on the following. What conclusions do you draw? Do the variances of the principal components add up to 4 (which is the number of variables)?

```
M <-structure(c(-1.4, -1.09, -0.19, -0.09, 1.04, -0.5, -0.22, -1.65, -1.02,
0.71, 0.23, -0.48, -1.4, -1.2, -0.5, 0.19, 0.9, -0.59, -0.23, -1.9, -0.73,
0.84, 0.21, -0.46, -1.62, -1.02, 0.03, 0.18, 1.11, -0.62, -0.12, -1.45, -1.13,
0.93, 0.4, -0.52, -1.52, -1.11, -0.24, -0.9, 0.99, -0.38, -0.04, -1.54, -1.15,
0.37, 0.48, -0.35),
.Dim = c(12L, 4L), .Dimnames = list(NULL, c("X1", "X2", "X3", "X4")))
```

### Solution:

Running the correlation-based PCA we get the following,

### Code:

```
> summary(princomp(M, cor = TRUE))
Importance of components:

               Comp.1      Comp.2      Comp.3      Comp.4
Standard deviation    1.9551976  0.36853202  0.194175456  0.0606830718
Proportion of Variance 0.9556994  0.03395396  0.009426027  0.0009206088
Cumulative Proportion 0.9556994  0.98965336  0.999079391  1.0000000000

> c(0.9556994*4, 0.03395396*4, 0.009426027*4, 0.0009206088*4)
[1] 3.822797600 0.135815840 0.037704108 0.003682435

> sum(c(0.9556994*4, 0.03395396*4, 0.009426027*4, 0.0009206088*4))
[1] 4
```

Looking at the proportion of variance explained it seems clear the the data are highly one dimensional since we used the correlation based approach we can be sure that this behavior is not due to variable scaling. I would likely conclude that the data could be successfully reduced to a single dimension with a very little drop in variance.

**Exercise 4:** Find any multivariate data set of interest to you and perform PCA. The data should have at least three columns. Did PCA help to reduce dimensionality? Looking at the loadings of PC1, what does the PC seem to represent? Try to find a data set that (as far as you can tell) has not already had PCA performed on it online?

### Solution:

The data set I used comes from a survey of couples on the subject of their relationship

with the goal building a model to predict if the couple is divorced or not. The data set has 54 predictors which represent survey questions with participants responding on an integer scale from 0 to 4. The data was downloaded from UC Irvine open source machine learning data repository. Performing a PCA(covariance based) using `princomp()` we find that almost 77% of the variance is explained by a single component. Looking at the loadings for that component it seems to represent a fairly even weighted sum, with only survey questions 6 and 7 contributing noticeably less than the others.

**Code:**

```
> divorce <- read_delim("divorce.csv", delim = ";",
+   escape_double = FALSE, trim_ws = TRUE)
```

```
> head(divorce)
  Atr1  Atr2  Atr3  Atr4  Atr5  Atr6  Atr7  Atr8  Atr9  ...
<dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1     2     2     4     1     0     0     0     0     0
2     4     4     4     4     4     0     0     4     4
3     2     2     2     2     1     3     2     1     1
4     3     2     3     2     3     3     3     3     3
5     2     2     1     1     1     1     0     0     0
6     0     0     1     0     0     2     0     0     0
```

```
> PCA <- princomp(divorce[, (1:54)])
```

```
> summary(PCA)
```

Importance of components:

	Comp.1	Comp.2	Comp.3	Comp.4	Comp.5
Comp.6					
Standard deviation	9.9135049	2.11850488	1.80288198	1.37424785	1.30018960
Proportion of Variance	0.7722181	0.03526505	0.02553995	0.01483938	0.01328308
Cumulative Proportion	0.7722181	0.80748315	0.83302310	0.84786248	0.86114556
	Comp.7	Comp.8	Comp.9	Comp.10	
Comp.11					
Comp.12					
Standard deviation	1.19411314	1.096292351	1.066951587	1.044745300	0.974016988
Proportion of Variance	0.01120409	0.009443615	0.008944889	0.008576427	0.007454501
Cumulative Proportion	0.88419830	0.893641917	0.902586806	0.911163233	0.918617734
	Comp.13	Comp.14	Comp.15	Comp.16	
Comp.17					
Comp.18					
...					
Standard deviation	0.886145774	0.831871787	0.805137992	0.769054831	0.762101073
Proportion of Variance	0.006170152	0.005437488	0.005093616	0.004647294	0.004563633
Cumulative Proportion	0.931764066	0.937201554	0.942295169	0.946942463	0.951506096

```
> PCA$loadings[,1]
```

	Atr1	Atr2	Atr3	Atr4	Atr5	Atr6	Atr7
Atr8	0.14532355	0.12780166	0.11695844	0.12890181	0.15604871	0.02763893	0.04136702
	Atr9	Atr10	Atr11	Atr12	Atr13	Atr14	Atr15
Atr16							


```

0.14766593 0.12400143 0.15605523 0.13299777 0.13456220 0.13602599 0.13901268 0.13681112
  Atr17      Atr18      Atr19      Atr20      Atr21      Atr22      Atr23
Atr24
0.15471569 0.14788853 0.15783402 0.14721578 0.13656669 0.13260662 0.15059101 0.13711374
  Atr25      Atr26      Atr27      Atr28      Atr29      Atr30      Atr31
Atr32
0.14097043 0.13749577 0.13598832 0.13245854 0.15275977 0.13880466 0.14034432 0.14282389
  Atr33      Atr34      Atr35      Atr36      Atr37      Atr38      Atr39
Atr40
0.16736545 0.14678003 0.17401198 0.17115496 0.15714922 0.16058318 0.15960030 0.16923167
  Atr41      Atr42      Atr43      Atr44      Atr45      Atr46      Atr47
Atr48
0.16133608 0.12565843 0.08606940 0.14735120 0.09421994 0.07399709 0.11971012 0.07918526
  Atr49      Atr50      Atr51      Atr52      Atr53      Atr54
0.12011115 0.11596136 0.09802381 0.10463189 0.11631775 0.14813410

> barplot(PCA$loadings[,1])

```

Figure 4: Bar plot of Loadings for each question in the survey.



Looking at the scores of the observations in the first component we see they exhibit clear bimodality, where the boundary seems to designate the divorced and non-divorced observations (around observation 85).

**Code:**

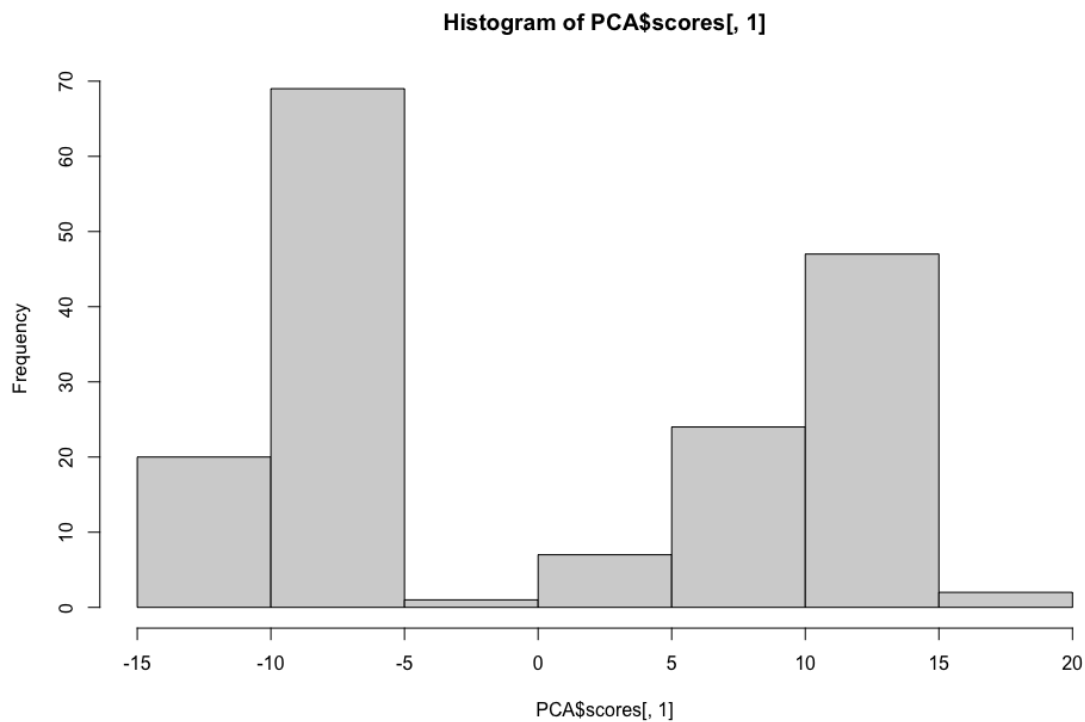
```

> PCA$scores[,1]
[1] -5.034834  6.082489  2.112230  4.765050 -6.048864 -6.125519
[7]  4.636226  2.866737  1.245790 -3.994374 15.013718 15.013718
[13] 13.684812 13.684812 13.537460 13.682746 13.682746 13.582610
[19] 13.140429 13.582610 13.116951 13.116951 12.835676 12.464237
[25] 12.551155 12.464237 12.391429 12.370017 12.375253 12.462171
[31] 12.375253 12.362035 12.340623 12.259620 12.275117 11.809457

```

[37]	11.867411	11.634894	11.634894	11.685843	11.149442	11.241596
[43]	11.281387	11.154678	11.241596	11.133265	11.154678	11.133265
[49]	11.241596	11.154678	11.241596	11.060458	11.009651	10.715591
[55]	10.936844	10.692113	10.505059	10.505059	10.378350	9.934103
[61]	9.934103	9.912690	9.934103	9.934103	9.960676	9.723570
[67]	9.702158	9.789076	9.629350	8.642267	9.406032	8.500951
[73]	8.629705	8.255912	8.059799	7.111691	5.882899	5.454844
[79]	5.650616	5.044036	4.252657	9.973884	7.241175	2.845738
[85]	-11.725144	-11.659492	-11.261800	-11.376719	-10.998180	-10.725802
[91]	-11.025546	-10.991772	-10.130046	-10.071233	-9.919028	-10.279158
[97]	-10.105848	-10.118556	-10.541153	-10.041946	-10.130887	-10.133866
[103]	-10.147190	-9.969850	-10.152863	-9.652144	-10.127179	-9.651935
[109]	-9.552364	-9.510380	-9.488157	-9.383214	-9.519640	-9.028921
[115]	-9.080701	-9.678057	-9.672411	-9.054363	-9.425221	-9.493400
[121]	-9.396816	-8.885112	-9.122728	-9.453767	-9.472082	-8.762338
[127]	-8.762338	-9.297598	-8.279844	-9.216059	-8.830888	-8.791266
[133]	-9.323096	-8.888367	-9.008702	-9.027225	-8.669607	-9.186244
[139]	-9.186244	-8.736300	-8.841863	-8.533205	-8.887203	-8.802613
[145]	-8.689610	-9.051975	-8.557702	-8.789391	-8.568272	-8.918316
[151]	-8.917761	-8.381054	-8.689880	-8.442909	-8.197338	-8.298847
[157]	-8.291527	-8.337184	-8.112729	-8.214733	-8.122544	-8.010257
[163]	-8.883499	-8.182890	-8.299792	-7.834202	-8.253722	-8.121778
[169]	-8.167597	-8.319674				

Figure 5: Histogram of Scores in the first component for each participant in the survey.



Not only can PCA reduce the dimensionality of the data significantly, but the loadings of the first component already give us a good model to classify the data (when we consider the scores). The loadings also tell us which survey questions we might want to question the efficacy of (questions 6 and 7).